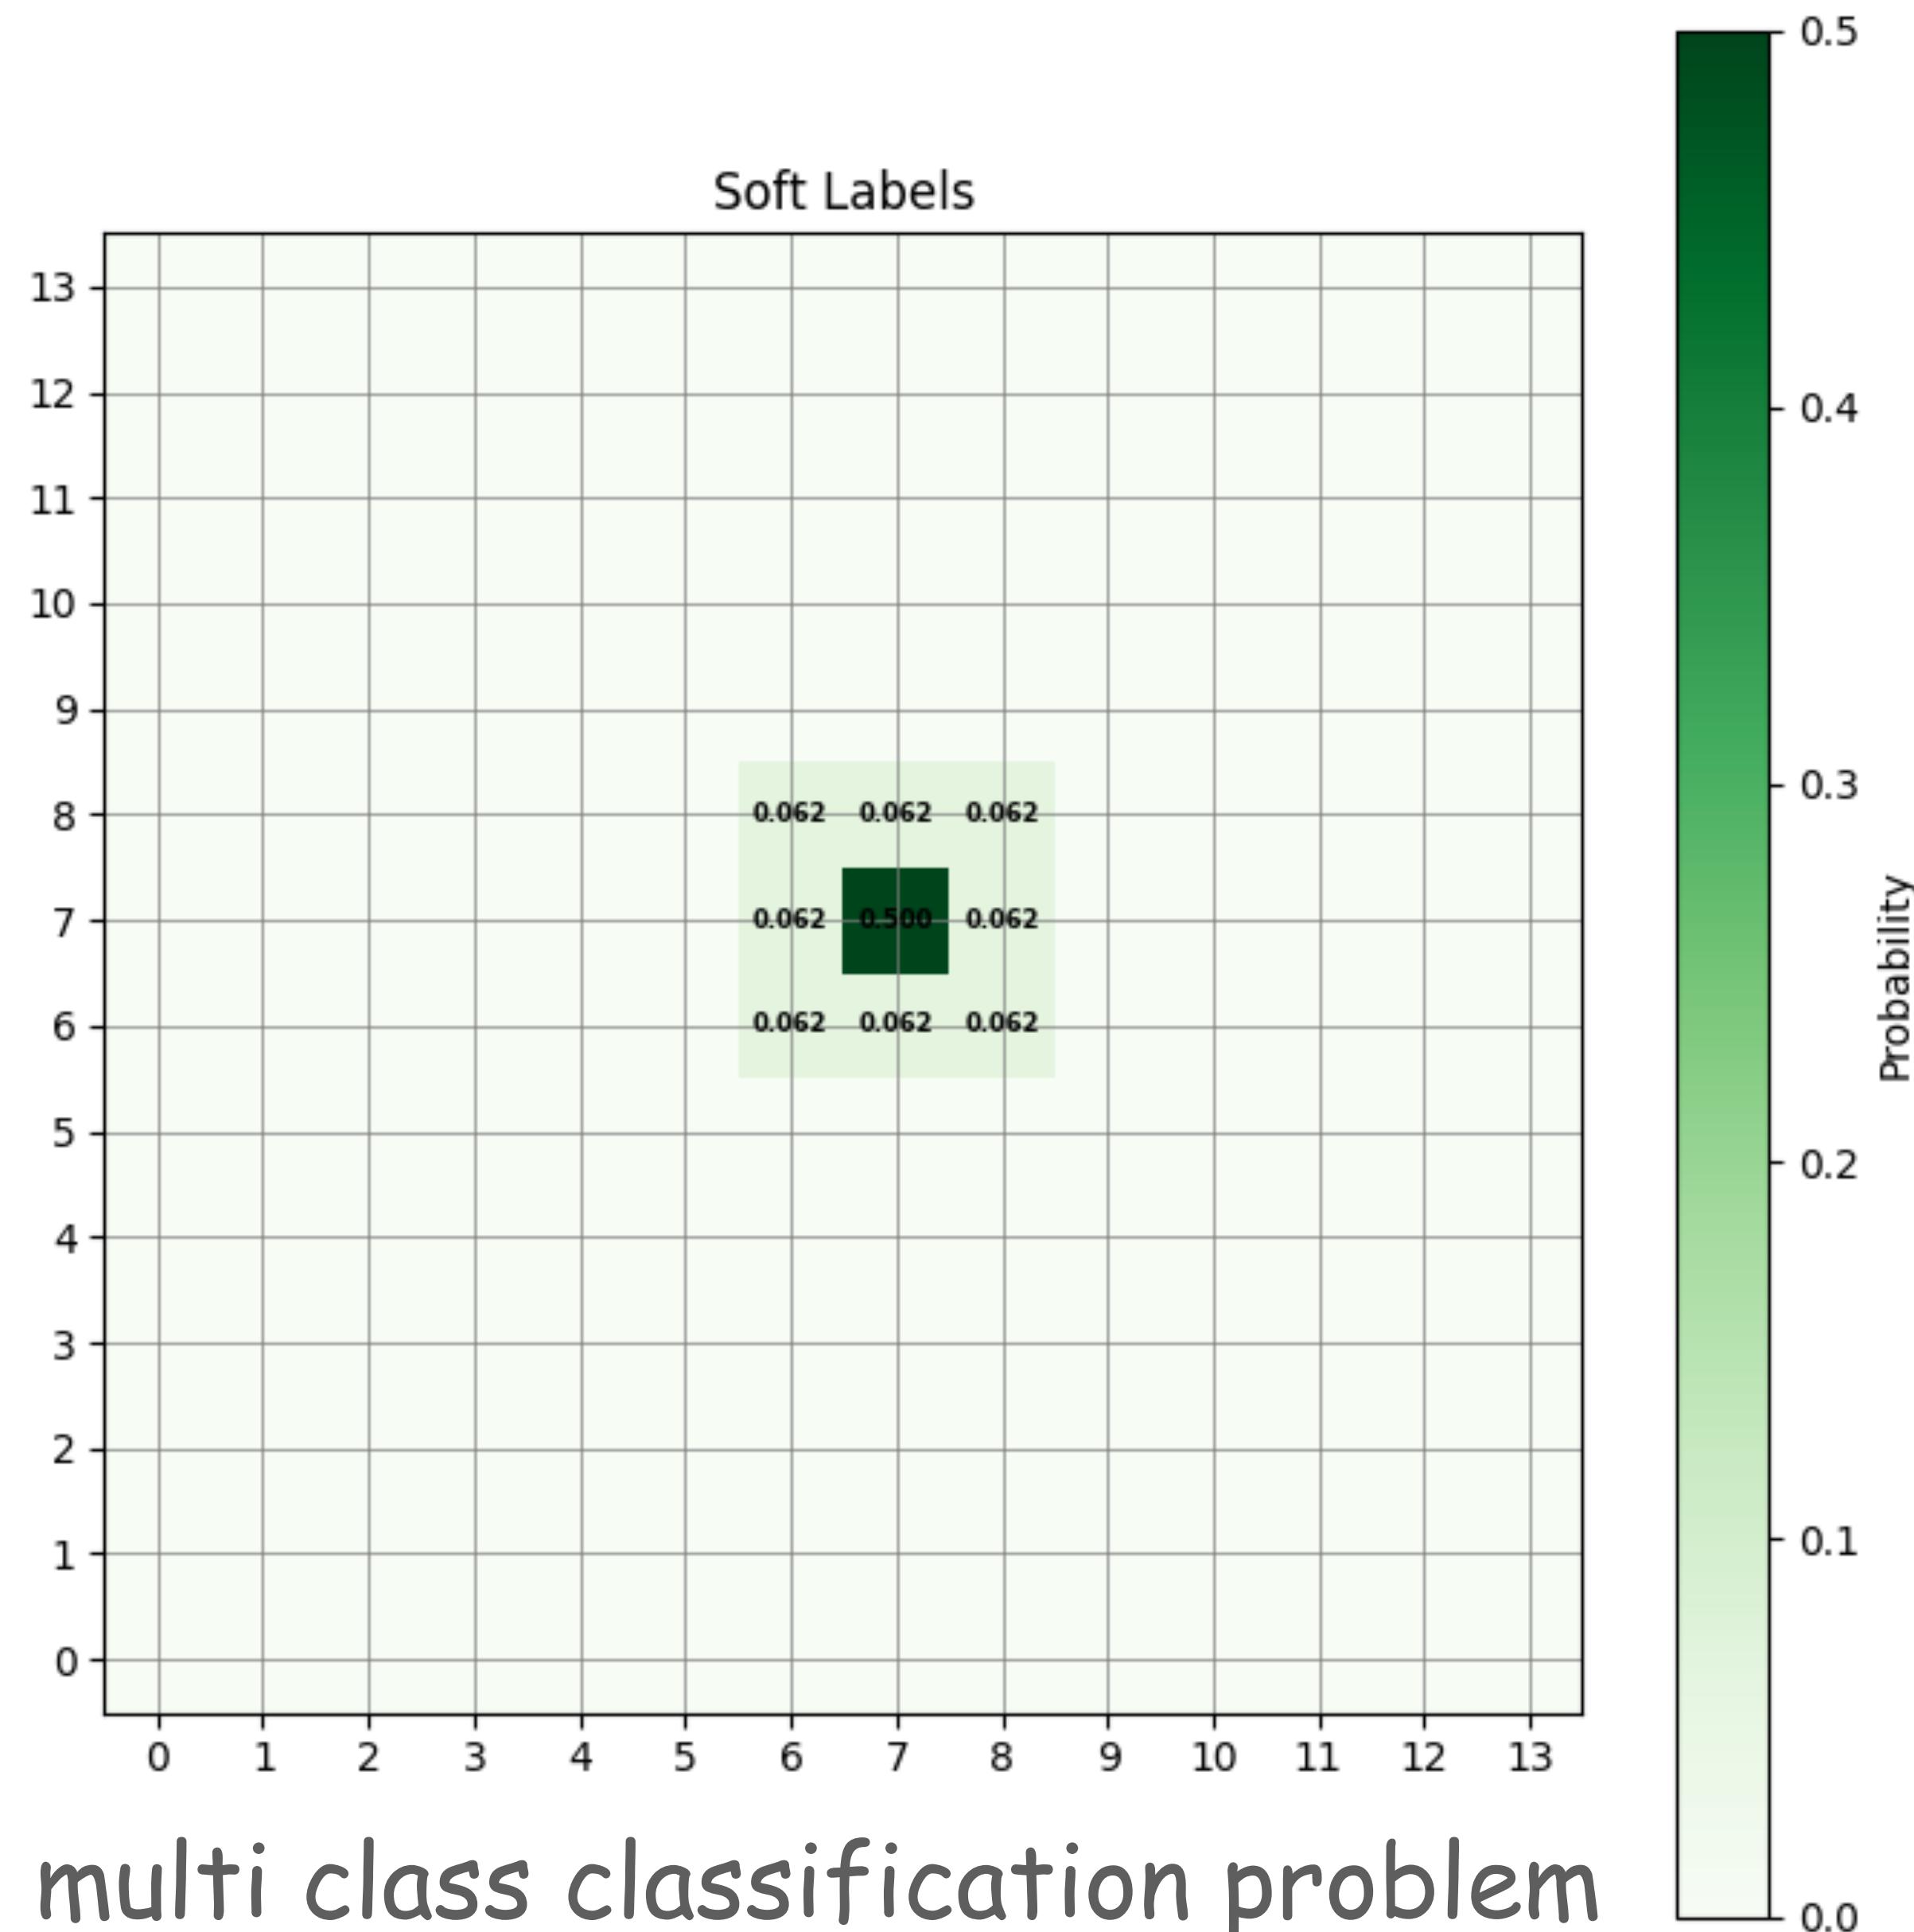


Keypoint Loss : Ground Truth soft labelling



Geometric Priors : Distance Loss

$$\mathcal{L}_{\text{dist}} = \frac{1}{|\mathcal{V}|} \sum_{(n,k) \in \mathcal{V}} H_\beta(\mathbf{p}_{nk} - \mathbf{g}_{nk})$$

Penalizes per keypoint deviation between prediction and ground truth